Homework #7

For the following problems: CartPole, MountainCar and LunarLander, implement the following strategies for optimal control using linear function approximation.

1) Monte Carlo
2) Semi-gradient SARSA
3) Semi-gradient QLearning

You need to decide on what you will use for $x(s)$, but at least you should use one of:

- State aggregation
- RBF, polynomial or Fourier basis
- Tiling

For each one of the problems, you need to turn in:

- Source code, explanation of the encoding function $x(s)$ for each problem
- Short movie playing the environment
- Graph of training rewards over episodes played